

# untitled199

September 8, 2024

```
[1]: import numpy as np
```

```
[2]: import pandas as pd
```

```
[3]: import matplotlib.pyplot as plt
```

```
[28]: import seaborn as sns
```

```
[5]: from sklearn.cluster import KMeans  
from sklearn.preprocessing import StandardScaler
```

```
[7]: data =pd.read_csv("mall_customers.csv")
```

```
[8]: data
```

```
[8]:
```

	customer_id	gender	age	annual_income	spending_score
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40
..	...	...	...	...	...
195	196	Female	35	120	79
196	197	Female	45	126	28
197	198	Male	32	126	74
198	199	Male	32	137	18
199	200	Male	30	137	83

[200 rows x 5 columns]

```
[10]: print(data.head(10))
```

	customer_id	gender	age	annual_income	spending_score
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

5	6	Female	22	17	76
6	7	Female	35	18	6
7	8	Female	23	18	94
8	9	Male	64	19	3
9	10	Female	30	19	72

```
[11]: data.describe()
```

```
[11]:
```

	customer_id	age	annual_income	spending_score
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[12]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   customer_id     200 non-null   int64
1   gender          200 non-null   object
2   age             200 non-null   int64
3   annual_income   200 non-null   int64
4   spending_score  200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[21]: # Print the column names to check for discrepancies
print(data.columns)
```

```
Index(['customer_id', 'gender', 'age', 'annual_income', 'spending_score'],
      dtype='object')
```

```
[23]: # Select features for clustering
X = data[['annual_income', 'spending_score']]

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print(X_scaled[:5]) # Preview the scaled features
```

```
[[-1.73899919 -0.43480148]
 [-1.73899919  1.19570407]
 [-1.70082976 -1.71591298]
 [-1.70082976  1.04041783]
 [-1.66266033 -0.39597992]]
```

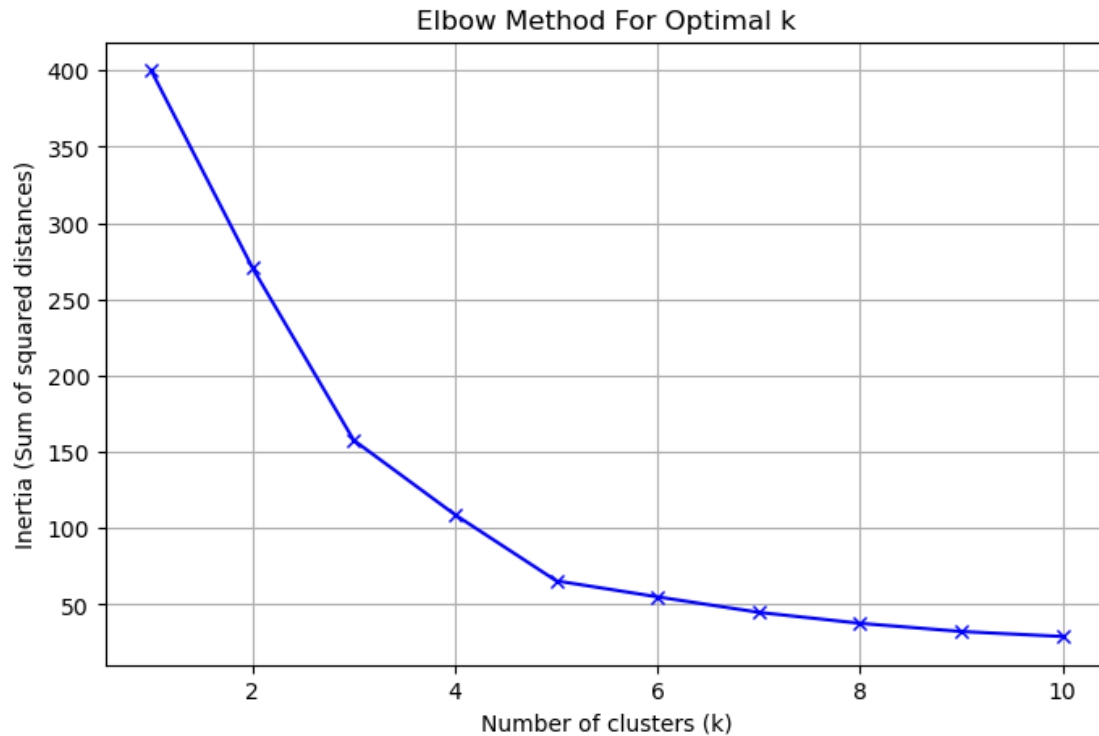
```
[25]: # Elbow Method to find the optimal number of clusters
```

```
inertia = []
K = range(1, 11)

for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)

# Plot the elbow curve
plt.figure(figsize=(8, 5))
plt.plot(K, inertia, 'bx-')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Inertia (Sum of squared distances)')
plt.title('Elbow Method For Optimal k')
plt.grid(True)
plt.show()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1036:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
  warnings.warn(
```



```
[26]: # Apply K-Means with optimal clusters (e.g., k=5)
kmeans = KMeans(n_clusters=5, random_state=42)
kmeans.fit(X_scaled)

# Get the cluster labels and add them to the original dataset
data['Cluster'] = kmeans.labels_

# Print centroids
centroids = kmeans.cluster_centers_
print(f"Centroids:\n {centroids}")
```

```
Centroids:
[[-0.20091257 -0.02645617]
 [ 1.05500302 -1.28443907]
 [-1.30751869 -1.13696536]
 [-1.32954532  1.13217788]
 [ 0.99158305  1.23950275]]
```

```
[30]: # Visualizing the clusters
plt.figure(figsize=(10, 7))

# Scatter plot for each cluster
```


```

sns.scatterplot(x='annual_income', y='spending_score', hue='Cluster',
               ↪data=data, palette='Set1', s=100)

# Plot the centroids
plt.scatter(centroids[:, 0] * scaler.scale_[0] + scaler.mean_[0], #
           ↪Denormalize the centroids
           centroids[:, 1] * scaler.scale_[1] + scaler.mean_[1],
           s=300, c='black', marker='X', label='Centroids')

plt.title('Customer Segments based on Income and Spending Score')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
plt.grid(True)
plt.show()

```



```

[31]: # Group the data by cluster and calculate the mean of each feature
cluster_summary = data.groupby('Cluster').mean()
print(cluster_summary)

```

```

customer_id    age  annual_income  spending_score

```

Cluster

0	86.320988	42.716049	55.296296	49.518519
1	164.371429	41.114286	88.200000	17.114286
2	23.000000	45.217391	26.304348	20.913043
3	23.090909	25.272727	25.727273	79.363636
4	162.000000	32.692308	86.538462	82.128205

[ ]: